

A Comprehensive Performance Analysis of Various Information Retrieval Algorithms for TRECCAR Problem

Internal documentation

Sumanta Kashyapi · Dr. Laura Dietz

the date of receipt and acceptance should be inserted later

Abstract A large set of randomly chosen Wikipedia articles are modified into a well-documented corpus. TRECCAR problem is to assign paragraphs contained in each article in that corpus to its appropriate sections. This paper consists of implementation details of three major information retrieval algorithms; LDA topic modeling, KMeans and Unigram topic modeling and their performance analysis over the TRECCAR problem. An evaluation framework is developed to measure performance regarding quality of clusterings and assignments. Obtained results are compared based on various measures calculated by the evaluation framework.

Keywords Topic modeling · LDA · Information retrieval · KMeans clustering · Unigram topic model · TRECCAR

1 Introduction

Categorization of large set of textual materials is a well traversed track of research in the field of information retrieval. Typically in these problems we have a collection of text documents as our input. Using the textual information contained in each document we are to determine to which pre-defined category it belongs. In our case, we have a set of M articles ($A_m \mid m \in \mathbb{N}, m \leq M$). For each article A_m we have a set of K_m section labels ($S_{mk} \mid k \in \mathbb{N}, k \leq K_m$) and a set of N_m textual content or paragraphs ($P_{mn} \mid n \in \mathbb{N}, n \leq N_m$) relevant to the article. For each article we have to map its section labels to paragraphs based on the information we extract from the set of paragraphs. The assignment results for all the articles are compared with the correct assignment provided

Sumanta Kashyapi
University of New Hampshire
E-mail: sumantakashyapi@gmail.com

Dr. Laura Dietz
University of New Hampshire

in a ground truth file. An evaluation framework is developed to quantify this comparison which in turn can be utilized to determine which categorization are better than the rest. Various categorization and classification algorithms can be used for the assignment task and for each approach we can measure it's effectiveness using the evaluation framework. For the experiments described here, we have used LDA topic modeling, KMeans, Unigram topic modeling and some of their variations. Our aim will be to identify which components of these algorithms are responsible for good results from the performance measures calculated by our evaluation framework. We will also carry out a feasibility study of a hybrid approach consisting of the best components for this task.

2 LDA topic modeling

Latent Dirichlet Allocation or LDA is a generative probabilistic model whose application in the field of information retrieval (specially in NLP) is pioneered by Dr. David Blei with his LDA topic model approach. This approach assumes that every document follows a random mixture of latent topics where each topic is a distribution over the vocabulary of the corpus. Documents are created by a generative process described by algorithm . This means if we

```

for  $k = 1$  to  $K$  do
  |  $\phi^{(k)} \sim \text{Dirichlet}(\beta)$ 
end
foreach document  $d \in D$  do
  |  $\theta_d \sim \text{Dirichlet}(\alpha)$ 
  | foreach word  $w_i \in d$  do
  | |  $z_i \sim \theta_d$ 
  | |  $w_i \sim \phi^{(z_i)}$ 
  | end
end

```

Algorithm 1: LDA generative process

can learn θ, ϕ and z from our given data w and a reasonable guess of α and β then we could know the most probable words that associates it's underlying latent topics. This means we would have an idea of the latent topics their respective proportion in a document which generated the document in the first place. Formally we want to reverse the generative process and want to learn the posterior distribution of the latent variables which is as follows:

$$p(\theta, \phi, z \mid w, \alpha, \beta) = \frac{p(\theta, \phi, z, w \mid \alpha, \beta)}{p(w \mid \alpha, \beta)} \quad (1)$$

As this equation is computationally intractable, other approximation strategies are used to get close to the actual values. Gibbs sampling method proved to be very effective in this process. It is a member from the Markov Chain Monte Carlo (MCMC) set of algorithms. The LDA equation that we will solve using

the Gibbs sampling is as follows (for more details on the derivation please refer):

$$p(z_i | z^{(-i)}, w) \propto (n_{d,k}^{(-i)} + \alpha_k) \frac{n_{k,w}^{(-i)} + \beta_w}{\sum_{w'} n_{k,w'}^{(-i)} + \beta_{w'}} \quad (2)$$

Gibbs sampling algorithm is implemented using the algorithm 2.

```

Initialize z
foreach iteration do
  for i = 0 to N - 1 do
    word ← w[i]
    topic ← z[i]
    nd,topic − = 1, nword,topic − = 1, ntopic − = 1
    for k = 0 to K - 1 do
      p(z = k | ·) = (nd,k + αk)  $\frac{n_{k,w} + \beta_w}{n_k + \beta \times W}$ 
    end
    samplw topic from p(z | ·)
    z[i] ← topic
    nd,topic + = 1, nword,topic + = 1, ntopic + = 1
  end
end
return z, nd,k, nk,w, nk

```

Algorithm 2: Gibbs sampling for LDA algorithm

2.1 Modifications and implementation details for assignment task

We took the basic implementation of LDA algorithm (cc.mallet.topics.SimpleLDA) from *mallet* API and customized it according to our requirements. For example, *SimpleLDA* implements *typeTopicCounts*, the data structure that stores the number of times one type/token is assigned to one topic, is implemented as a simple 2D integer array. However, the topic inferencer for LDA (cc.mallet.topics.TopicInferencer) implements the same data structure in a different way to save space. For each type/token the count of topic assignment is stored in a single integer object using bitshift operations. Therefore we had to implement the same strategy in *SimpleLDA* in order to make it compatible with *TopicInferencer*.

3 KMeans algorithm

KMeans is a simple but effective unsupervised clustering algorithm. It works by repeatedly calculating and placing centroids for each cluster until the centroids do not move. First, for each cluster is strategically placed in the feature space so that we have a better chance at achieving global optima. Then using some metric function, a cluster is assigned to every data points for which its centroid is nearest to the data point. After all the data points are assigned some

cluster, centroids are recalculated using the same metric function. This loop continues until centroids do not move and the algorithm converges. Formally, the algorithm tries to minimize the below objective function.

$$J = \sum_{j=1}^k \sum_{i=1}^n \|x_i^{(j)} - c_j\|^2 \quad (3)$$

where $\|x_i^{(j)} - c_j\|^2$ is the metric distance between cluster c_j and data element $x_i^{(j)}$. The basic algorithm that any typical KMeans implementation will follow is presented in algorithm 3. Initializing the centroids is an important step in the algorithm because based on the initial guess our final results and also the number of iterations needed to reach that result may vary.

```

Data: paragraph list, no. of clusters
Result: cluster of paragraphs
initialize centroids/cluster means  $m_1, m_2, \dots, m_k$ 
while changes in centroids do
  for each  $x_i$  where  $i = 1$  to  $n$  do
    | assign  $x_i$  to the cluster  $m_j$  having  $\min \|x_i^{(j)} - m_j\|^2$ 
  end
  for each  $m_j$  where  $j = 1$  to  $k$  do
    | Replace  $m_j$  with the mean of all of the  $x_i^{(j)}$  in cluster  $j$ 
  end
end

```

Algorithm 3: KMeans algorithm

3.1 Modifications and implementation details for assignment task

Using KMeans algorithm in our experiments proved to be quite straight forward. We simply imported the necessary packages (`cc.mallet.cluster`) from `mallet` API and called up necessary methods without any modifications. However, we encountered certain issues for which we had to modify our existing codes for topic model to accomodate KMeans. Topic models in `mallet` are implemented in such a way that it internally converts the data elements (Instance objects) into a sequence of index, value pairs called *SparseVector*. Whereas in case of KMeans, it works with sequence of feature counts or *FeatureSequence*. This conversion along with all the pre-processing tasks are carried out by a chain of processing elements known as pipes (`cc.mallet.pipe`) and are placed inside `buildPipe()` method of *SingleRun* class. To take account for the different input format required by KMeans algorithm, we had to modify `buildPipe()` so that it can provide pre-processed data item in proper format. It is to be noted that, this difference is only related to representation of data items which do not interfere with actual data values and hence results obtained from both class of algorithms remains comparable.

Although we used almost unmodified version of `mallet KMeans` algorithm for clustering the paragraphs, we still had to develop our own strategy to assign sections to these clustered paragraphs. The algorithm is implemented in `assignUsingKMeans()` to take care of the assignment task.

```

Data: kmeans object, paragraph list, section ID list
Result: section ID to list of para IDs mapping
foreach section sec in list of sections do
    | foreach cluster mean cm in list of cluster means do
    | | secMatrix[sec][cm] = metric(sec, cm)
    | end
end
foreach paragraph p in list of paragraphs do
    | foreach cluster mean cm in list of cluster means do
    | | paraMatrix[p][cm] = metric(p, cm)
    | end
end
call matrixAssignment() with secMatrix and paraMatrix

```

Algorithm 4: Assignment algorithm for KMeans

4 Unigram topic model and it's implementation

In his paper, author Blei describes how his LDA model extended from the simple Unigram mixture model. According to the unigram mixture model the probability of a document $p(\mathbf{w})$ in a corpus can be expressed as:

$$p(\mathbf{w}) = \sum_z p(z) \prod_{n=1}^N p(w_n | z)$$

However, in this paper we considered modifying the existing LDA implementation to have an additional constraint that will enforce unigram behaviour to typical LDA model. The constraint we tried to implement here is each token in a single document has to agree upon a single topic in a single iteration. So we simply would not allow a state during any iterations in which some token is assigned a topic different than it's adjacent ones. We did this by assigning topics to an entire document rather than individual tokens. However, we still considered individual tokens for probability calculations of the next topic assignment. Calculation of the distribution $p'(m)$, from where \tilde{k} is drawn, dominates the training phase. According to unigram model a topic is assigned to a whole document m and therefore all tokens $t = 1$ to T_m inside that document m takes the same topic \tilde{k} or $z_t = \tilde{k}$ for all $t = 1$ to T_m . Hence each probability in the distribution $p'(m)$ can be expressed as:

$$p'(m_i | \vec{m}_{-i}, \vec{w}) = \prod_{t=1}^{T_m} p_t(z_t = k | \vec{z}_{-t}, \vec{w})$$

Data: Pre-processed corpus
Result: Trained Unigram topic model for input corpus
for each doc m in corpus where $m \in [1, M]$ **do**
 for each topic $i=1$ to k **do**
 | calculate $p'(m_i | \vec{m}_{-i}, \vec{w})$;
 end
 sample a new topic $\tilde{k} \sim p'(m)$;
 for each token $t \in [1, T_m]$ in doc m **do**
 | decrement count from all data structures involving old topic assignment;
 | assign new topic \tilde{k} to all data structures;
 end
end

Algorithm 5: Unigram topic model training phase

Being an extension of original LDA algorithm, we can still use the proportional relation

$$p_t(z_t = k | \vec{z}_{-t}, \vec{w}) \propto \frac{n_{k,-t}^{(v)} + \beta_v}{\sum_{v=1}^V (n_{k,-t}^{(v)} + \beta_v)} (n_{m,-t}^{(k)} + \alpha_k)$$

But $n_{m,-t}^{(k)} = 0$ for all m and k , so our proportional equation simplifies to:

$$\begin{aligned} p'(m_i | \vec{m}_{-i}, \vec{w}) &\propto \prod_{t=1}^{T_m} \frac{n_{k,-t}^{(v)} + \beta_v}{\sum_{v=1}^V (n_{k,-t}^{(v)} + \beta_v)} \\ &\text{ignoring proportionality constant and assuming symmetrical } \beta \text{ distribution over vocabulary} \\ &= \prod_{t=1}^{T_m} \frac{n_{k,-t}^{(v)} + \beta}{\sum_{v=1}^V n_{k,-t}^{(v)} + \beta \text{Sum}} \\ &= \exp \left[\sum_{t=1}^{T_m} \log \left(n_{k,-t}^{(v)} + \beta \right) - T_m \log \left(\sum_{v=1}^V n_{k,-t}^{(v)} + \beta \text{Sum} \right) \right] \\ &\text{in log space} \end{aligned}$$

While implementing unigram model we computed $p'(m_i)$ in log space to avoid underflows. Also instead of directly sampling from p' , we used the smoothed version (Jelinek-Mercer smoothing) p'' which is calculated as:

$$p''(m_i) = \lambda p'(m_i) + (1 - \lambda)$$

Here λ is inverse temperature which controls randomness during sampling. In other words, low λ leads to more *exploration* from wide range of topics and high λ will *exploit* from narrow range of topics. The topic inferencer for unigram model, which is used during the section to paragraph assignment, performs similar calculations except, instead of paragraphs in corpus it takes the text content related to section headings.

4.1 Unigram topic inferencer

We followed algorithm 6 while developing Unigram topic inferencer (... playground.topics.UnigramTopicInferencer). After developing the training model for Unigram topic model, we had to develop an topic inferencer which would be compatible with existing evaluation framework. The method signature had to be same as the LDA inferencer but internally it had to calculate topic probabilities using Unigram model.

```

Data: Section data
Result: Inferred topic scores
for each topic  $k=1$  to  $numTopics$  do
    for each token  $t=1$  to  $numTokens$  do
        calculate  $sumLog_k += \log \left( n_{k,-t}^{(v)} + \beta \right)$ ;
        /*  $n_{k,-t}^{(v)}$  is represented by  $typeTopicCounts[v][k]$  where  $v$  is the type
           retrieved from vocabulary at token position  $t$  */
    end
    calculate
     $topicScores[k] = \exp \left[ sumLog_k - numTokens \log \left( \sum_{v=1}^V n_{k,-t}^{(v)} + \beta Sum \right) \right]$ ;
    /*  $\sum_{v=1}^V n_{k,-t}^{(v)}$  is implemented in  $tokensPerTopic[k]$  */
end
return normalized  $topicScores$ ;

```

Algorithm 6: Unigram topic inferencer algorithm

5 Evaluation Framework

Before we try to experiment with different combinations of topic modeling and clustering modules with our treccar corpus, we must have a reliable framework to compare our results. A robust evaluation framework for our experiments is a key element which will help us to be confident about our findings. We will evaluate two aspects of our results. First, we have a set of measures which will determine the accuracy of the clustering that we have done with our set of data elements, which is in this case a set of paragraph objects identified by unique paragraph ids. Note that while measuring how good our clusters are compared to that of the ground truth, we do not utilize the assignment mappings or more precisely the section objects under which each cluster is mapped. In our second set of measure we do just that by comparing our section to paragraph mappings with that of ground truth.

5.1 Clustering measures

In our current evaluation framework we are using two basic statistical measures for evaluating our clustering results.

Adjusted RAND index RAND index is a popular statistical measure used to quantify the similarity between two clusterings done on the same dataset. In our case those two clusterings are resulting clusters from our designed model and the clusters made according to the ground truth file. Although RAND index is often associated with the accuracy of the results but in our case we are using it without the class labels. Typically RAND index, R is calculated as:

$$R = \frac{a + b}{a + b + c + d} = \frac{a + b}{\binom{n}{2}} \quad (4)$$

where a, b = pairs of paragraphs present in either same or different cluster for both clustering,

c, d = pairs of paragraphs present in same cluster in one but in different cluster in other,

n = no. of paragraphs But this simple RAND index expects the two partitions to have same number of clusters. However, we want the framework to be robust enough so that we can compare our results even when our resulting partition is of different size from that of the ground truth. Hence adjusted RAND index or ARI is used which is calculated as:

$$ARI = \frac{\sum_{ij} \binom{n_{ij}}{2} - [\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2}] / \binom{n}{2}}{1/2[\sum_i \binom{a_i}{2} + \sum_j \binom{b_j}{2}] - [\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2}] / \binom{n}{2}} \quad (5)$$

where n_{ij} = contingency table values from i th row and j th column

a_i = contingency table i th row sum

b_j = contingency table j th column sum

Purity Purity is a simple clustering measure which is used as an external evaluation criterion to measure quality of the clustering. While calculating purity, we basically measure the accuracy of the assignment if we assign each cluster to the most frequent class in that cluster. Hence we a high purity means that most of the data elements inside each cluster agrees upon the class labels. Purity measure is calculated using the equation 6.

$$Purity(M, C) = \frac{1}{N} \sum_k \max_j |m_k \cap c_j| \quad (6)$$

where $M = \{m_1, m_2, m_3, \dots, m_k\}$ set of clusters

$C = \{c_1, c_2, c_3, \dots, c_J\}$ set of classes according to ground truth

N = total number of data elements (paragraphs)

Table 1 TREC eval measures

Measure	Purpose
num_ret	Total number of documents retrieved over all queries
num_rel	Total number of relevant documents over all queries
num_rel_ret	Total number of relevant documents retrieved over all queries
map	Mean Average Precision (MAP)
gm_ap	Average Precision. Geometric Mean, $q_score = \log(MAX(map, .00001))$
R-prec	R-Precision (Precision after R (= num-rel for topic) documents retrieved)
bpref	Binary Preference, top R judged nonrel
recip_rank	Reciprical rank of top relevant document
ircl_prn.[num]	Interpolated Recall - Precision Averages at [num] recall
P[d]	Precision after [d] docs retrieved

5.2 Assignment measures

After we finish clustering the paragraphs and assign them with their respective cluster labels, we have to assign them to sections. The quality of the section to paragraph assignment is measured using an existing framework, TREC eval, which is designed specifically for this set of problems where we have to evaluate rankings of retrieved information.

Different measures of TREC eval Table presents a quick summary of all the measures calculated by the TREC eval tool. We will mainly consider num_rel_ret, MAP, Rprec and bpref measures for our experiments.

6 Implementation details

The algorithms, models and evaluation framework are developed in java. Mallet API has been used for basic text retrieval tasks and as a starting point for developing all the models used in the experiments.

6.1 Important classes

RunExperiment: It is the starting point of the evaluation framework. It contains most of the variables that control the workflow of the experiment including which algorithm and which version of it is to be used. It takes all the values for one set of experiments and constructs the topmost loop that calls the appropriate algorithm with it's respective values. Details of the variables are in Table

SingleRun: At each iteration of *RunExperiment*, *SingleRun* is called with appropriate arguments. This is the class which decides which version of which algorithm is to be called. Depending on the *RUN_BY_PAGE* varibale of *RunExperiment* class, *runExperiment()* or *runExperimentWholeCorpus()*

Table 2 RunExperiment variables

Variable	Purpose
SAVE_RESULT	Whether we want to save our results
RUN_BY_PAGE	Whether we want to run by page or take the whole corpus as a single page. In case it is false, the problem changes from section to paragraph assignment to page to paragraph assignment. In other words, instead of calculating measures for each page, we now replace all the section titles with page titles and we try to assign paragraphs to it's respective pages. Consequently we calculate a single measure for the entire corpus.
CLUSTERING_MEASURE_FILENAME	The name of the file which will store the clustering measures.
TRECEVAL_ASSIGN_FILENAME	The name of the file which will store the section to paragraph assignments in trec_eval format which will be the input for trec_eval script.
SMOOTHED_UMM	Whether we are using smoothing in Unigram topic model.
model	Which algorithm/model is to be used. Currently we have five valid strings for this: 1 = <i>LDA</i> , 2 = <i>KMeans</i> , 3 = <i>Unigram</i> , 98 = <i>Correct</i> , 99 = <i>Random</i>
[meta][Param]	These are the parameters for algorithms that we can vary. There are currently four parameters available which are k value, number of iterations for a single call to the respective algorithm, αSum and βSum two hyperparameters for topic modeling. Also for each parameter, we have three bounding values <i>[meta]</i> to calculate different values of the parameter throughout the iterations which are start, stop and step. For a particular parameter, we start off the iterations with start value of the parameter until we get to the stop value adding the step value to the current value after each iteration.
isVar	Each boolean value in this array represents one variable. Currently it is of size four and from 0 to 3 it represents $k, iterations, \alpha Sum$ and βSum respectively. For example, if we have the array as <i>[true, false, false, true]</i> then it means we will consider <i>stop</i> and <i>step</i> values of k and βSum and iterate through different values of k and βSum .

method of *SingleRun* is called from *RunExperiment*. Some important method details are discussed in Table.

ResultForPage: This class is generated for each page and holds the clusters of paragraph ids (*paraClusters*) and the mapping of section ids to list of paragraph ids (*queryParaAssignment*).

Table 3 Important SingleRun methods

Method	Purpose
runExperiment	Performs experiment on per page basis. This method is implemented using algorithm 7.
runExperimentWholeCorpus	Performs experiment on the whole corpus.
modelAndAssign	Based on the <i>model</i> parameter passed from <i>RunExperiment</i> different algorithms are called. This method is based on algorithm 8.
convInsAssignToIDAssign	Coverts mappings from Section Instance object to list of Paragraph instances to mappings from Section id to list of Paragraph ids.
assignUsing[algo]	These methods takes the clusters of paragraphs generated by some algorithm and assigns them to section instances using [algo] algorithm (<i>assignUsingKMeans</i> uses KMeans). Ideally any changes made to these methods will not affect the results of clustering measures because section to paragraph assignment information is not utilized while calculating those measures. Algorithm 9 is implemented in this method.
matrixAssignment	This is called from <i>assignUsing[algo]</i> to do the assignment using the KL divergence matrix (for topic modeling algorithms) or the distance matrix (for clustering algorithms).

```

foreach paragraph p in list of paragraphs do
  foreach query q in list of queries do
    calculate paraQueryMatrix[p][q] = assignVal for p and q
    /* assignVal is the KL-div value for topic modeling algorithms and
       metric distance for clustering algorithms so that the value
       becomes proportional to the quality of p to q assignment */
  end
end
initialize isParaAssigned[] = false
foreach paragraph p in list of paragraphs do
  bestQueryForPara = q where
  paraQueryMatrix[p][q] = min paraQueryMatrix[p]
end
while checkIfDone() do
  /* checkIfDone() will return true when all values in paraQueryMatrix
     is 0 */
  get index (p, q) of minimum positive value in paraQueryMatrix
  assign para[p] → query[q]
  put -1 in place of row p and column q of paraQueryMatrix
  isParaAssigned[p] = true
end
foreach paragraph p in list of paragraphs do
  if isParaAssigned[p] = false then
    | assign para[p] → bestQueryForPara[q]
  end
end

```

Algorithm 9: *assignUsing[modelName]()* algorithm

```

foreach page pg in corpus do
  | obtain clusters from ResultForPage[pg]
end
foreach key in gtMap do
  /* gtMap is the mapping between pageID and paragraph clusters
    according to ground truth file */
  if WHOLE_CORPUS_MODE then
    | get the list of lists of paragraphs as gtClusters
  end
  else if key starts with pageID then
    | add gtMap[key] to gtClusters
  end
  candClusters = clusters[pg]
  initialize contMat[gtClusters[pg].size()][candClusters.size()]
  // Contingency Matrix

  foreach row r in contMat do
    foreach column c in contMat do
      correct = gtClusters[r]
      candidate = candClusters[c]
      match = no. of matches between correct and candidate
      contMat[r][c] = match
    end
  end
end

```

Algorithm 10: Algorithm to form contingency matrix

MeasureExperiment: This class calculates the cluatering measures for a single page. Currently it calculates two measures, adjusted RAND index (*calculateRANDPerPage*) and Purity (*calculatePurityPerPage*) from the contingency matrix formed using the algorithm 10. Later on in *SingleRun* these values are used to get the mean and standard error for the current experiment.

6.2 Brief sequence of execution

We start the evaluation from *RunExperiment* with necessary arguments. It contains the top loop that calls *SingleRun* for each combination of the arguments provided. Inside the *SingleRun* one complete execution happens inside it's *runExperiment()* method. First of all, data elements such as list of paragraphs, sections, ground truth mappings are converted to *InstanceLists* which is recognized by the *mallet* API. For each article in the corpus, *modelAndAssign()* method is called which in turn calls *assignUsing[model]* based on the model-chosen in *RunExperiment*. For each page, we obtain a *ResultForPage* object which contains clusters and section to paragraph mappings. These result objects are handled by *MeasureExperiment* and we get the final evaluation results which are stored for later analysis.




Fig. 1 Minimum no. of iterations of LDA topic modeling for RAND




Fig. 2 Minimum no. of iterations of LDA topic modeling for Purity




Fig. 3 Minimum no. of iterations of LDA topic modeling for numq measures




Fig. 4 Minimum no. of iterations of LDA topic modeling for MAP and Rprec

Table 4 Minimum no. of iterations

Algorithm	min. iterations for RAND	purity	numq measures	MAP	Rprec	chosen iteration
LDA	> 80	> 40	NE	NE	NE	100
Unigram	NE	NE	NE	NE	NE	100

NE = No Effect

7 Experimental results

7.1 Baseline experiments

Two sets of experiments are carried out to form a baseline for other algorithms. In first set, we randomly assigned paragraphs to sections for each article. For this random assignment we got mean RAND index to be 0.0032. When it is ensured that every section got at least one paragraph, it comes down to -0.0014. more to add here purity, map, rprec

7.2 Minimum no. of iterations

Three sets of experiments were carried out for each algorithm to get the minimum number of iterations needed of the respective algorithm to reach sufficiently close to the highest result it can achieve. Each set of experiments were done using 4 set of measures; Adjusted RAND index, Purity, four numq measures from trec_eval tool and precision measures. The plots of this study for LDA algorithm are shown in figures 3 to 6 and for Unigram topic model algorithm are shown in figures 7 to 11.




Fig. 5 Minimum no. of iterations of Unigram topic modeling for RAND




Fig. 6 Minimum no. of iterations of Unigram topic modeling for Purity




Fig. 7 Minimum no. of iterations of Unigram topic modeling for numq measures




Fig. 8 Minimum no. of iterations of Unigram topic modeling for Rprec




Fig. 9 Minimum no. of iterations of Unigram topic modeling for MAP

7.3 Hyper parameter learning

LDA algorithm has two hyperparameters, α and β . We have to learn the optimum value for both of these parameters for our particular corpus using the evaluation framework. For Unigram topic model we have to learn only for β . To get the optimum values of these variables we run the algorithm over a sufficiently large range of a parameter while keeping rest at a constant value. Then we measure various evaluation metrics and find out a narrow range of parameter values for which we are getting sufficiently good measures.

Optimum α Sum and β Sum for LDA Figures from 9 to 12 plots result obtained from the experiments to learn optimum α Sum and β Sum for LDA algorithm. From these charts we found optimum range for α Sum to be 0.8 to 1.2 and optimum range for β Sum to be 100 to 300. For subsequent experi-




Fig. 10 Learning optimum αSum of LDA for RAND with βSum , iterations fixed




Fig. 11 Learning optimum αSum of LDA for Purity with βSum , iterations fixed




Fig. 12 Learning optimum βSum of LDA for RAND with αSum , iterations fixed




Fig. 13 Learning optimum βSum of LDA for Purity with αSum , iterations fixed




Fig. 14 Learning optimum βSum of UMM for RAND with 100 iterations (long range)




Fig. 15 Learning optimum βSum of UMM for Purity with 100 iterations (long range)

ments with LDA algorithm, we chose αSum and βSum values to be 1.0 and 260.

Optimum βSum for Unigram topic model As we have figured out from the iterations experiments, adding smoothing factor to Unigram model did not affect any measures when compared to its no-smoothing counterpart. Hence, from here onwards we will only use no-smoothing version of the Unigram model. Figure to presents the RAND index and purity plots for range of βSum values for Unigram model. Based on these plots, we have chosen $\beta Sum = 260$ as the optimum value for subsequent experiments with Unigram model.




Fig. 16 Learning optimum βSum of UMM for RAND with 100 iterations (short range)




Fig. 17 Learning optimum βSum of UMM for Purity with 100 iterations (short range)




Fig. 18 Comparison of RAND index vs. k




Fig. 19 Comparison of purity vs. k




Fig. 20 Comparison of RAND index at true no. of sections




Fig. 21 Comparison of Purity at true no. of sections

7.4 Comparison of algorithms with varying k

After we have decided the optimum values of hyperparameters for all the algorithms, we compare the measures from different algorithms by varying k and also at $k = \text{true no. of sections}$. This will give us an idea about the effectiveness of each algorithm for the task.

Figure 7.4 to 7.4 presents the comparison of 3 different algorithms based on clustering measures. Figure 7.4 to 7.4 presents the same comparison based on treceval measures.

8 Analysis of results

In previous section we have presented results from three different implementation of clustering and assignment models which are LDA topic modeling,




Fig. 22 Comparison of MAP vs. k




Fig. 23 Comparison of MAP at true no. of sections




Fig. 24 Comparison of Rprec vs. k




Fig. 25 Comparison of Rprec at true no. of sections




Fig. 26 Comparison of bpref vs. k




Fig. 27 Comparison of bpref at true no. of sections

KMeans clustering and Unigram mixture model. Based on these results we can gain following insights in context of our specific experimental setup.

From RAND index comparisons in Figure 7.4 and 7.4 it becomes evident that in terms of cluster similarity with that from the ground truth, Unigram model performs better than the rest. Although the purity comparison shows that KMeans algorithm constantly achieves high purity for high k value but we have to remember that it only means that the clusters are more specific or pure and does not signify anything regarding the overall quality of the clustering.

While trying to map section IDs to paragraph clusters using KMeans algorithm, we found around 2.5% of the section IDs are not considered for treceval measures. Consequently we obtained num_q measure for KMeans as 1489 instead of 1526. The reason behind this is for those section IDs all the

words/tokens are removed during stop word removal phase of pre-processing. So while measuring their metric distances from cluster means, a NaN is returned. This does not happen with topic modeling because according to the implementation topic probabilities are influenced by tokens of an instance. If an instance has an empty token list then topic probabilities are not influenced at all which means each topic has same probability to be picked for that instance. As random assignment of topics like this is valid in topic modeling, we get those sections in our final results which is not the case with KMeans.

The iteration experiments done with Unigram model is interesting because according to the result (Figure 7.2) it is clear to us that smoothing over pure Unigram topic probabilities degrades it's performance. This explains the sudden jump of smoothed version of Unigram model from close to 0 RAND index to it's no-smoothing version when number of iterations reaches 11. The smoothed UMM is implemented in such a way that it starts with $\lambda = 0.0$ or maximum smoothing (minimum contribution from UMM) and after each equal number of iterations the smoothing factor λ increases 0.1 until it reaches 1.0 (pure UMM). For no. of iterations upto 10, λ could not reach 1.0 and we receive very low RAND index. For iterations more than 10, λ reaches 1.0 at the end and we get RAND measure close to that of it's unsmoothed version. Another point worth noting here is that pure UMM seems to converge only after it's second iteration.

Provided the true number of sections as the value of k , it is found that LDA topic models performs slightly better than other algorithms based on all treceval measures shown here (Figures). This also shows that performing well in clustering tasks (Unigram model) does not guarantee success in assignment tasks.

9 Conclusion and Future scope

The experimental setups described here are far from perfect. More thorough investigation is needed to figure out the poor performance of all algorithms in assignment tasks and the underlying reason behind the success of Unigram model in clustering tasks over LDA despite being it's predecessor. Other than that, there are so many interesting directions one can follow from the results we already obtained here. The experiments presented in this paper basically focuses on performance anylysis of various algorithms on a page-by-page basis of the corpus. The next step would be to separate the corpus based on how difficult it is for a specific algorithm to achieve a good measure. This kind of difficulty analysis may lead to key insights into suitability of a particular algorithm for assignment and clustering tasks. Instead of assigning paragraphs to sections for each page, some quick modification to the evaluation framework will allow us to measure performance of these algorithms while assigning paragraphs to pages. This may allow the models to have more margin for error

and hence would be easier to compare them. Also this approach will allow us to include section IDs as part of page IDs and reduce the chances of getting empty Feature sequence or Vectors which happened to be a problem we faced in our experiments.

References

1. Author, Article title, Journal, Volume, page numbers (year)
2. Author, Book title, page numbers. Publisher, place (year)